

# CS 252: Systems Programming Fall 2024 Project 3: Shell Interpreter

## Part 1: Monday, September 23 11:59pm

### Goals

The goal is to build a shell interpreter combining behavior from common shells like Bash and csh. Some skeleton code is provided.

### Deadlines

- Part 1: Monday, September 23 11:59pm
  - Part 2: Wednesday, October 2 11:59pm
  - Part 3: Friday, October 11 11:59pm
- All extra credit parts are due with the final submission.

### Testing

Much of the shell will be graded using automated tests.

- `./testall partN` (where  $N = 1, 2, 3$ )
- `./testall` with no arguments runs all tests.

### The Assignment - Part 1

#### 1. Login and Build:

- Log in to a CS department machine and run `git clone ~cs252/repos/$USER`.
- Build the shell by typing `make` at the prompt.

#### 2. Parsing and Executing Commands

##### • Scanner and Parser:

- Use Lex and Yacc to write a scanner and parser for the shell grammar. The `struct sint` represents a shell command. The `struct command` represents a list of simple commands (which can be nested). The `struct error` represents an error redirection.
- Look through the `Makefile` to see how the scanner and parser are then compiled.

##### • Accepting more complex commands:

- Modify `shell.l` and `shell.y` to support the grammar:

```
cmd [arg]* [ | cmd [arg]* ]* [ [> filename] [< filename] [2> filename]
[ >& filename] [>> filename] [>>& filename] ]* [&]
```

- Test with commands like `ls`, `ls -al`, `ls -al aaa bbb cc`, `ls -al aaa bbb cc > outfile`, etc.

#### 3. Executing Commands

##### • Single command process creation and execution:

- For each single command, create a new process using `fork()` and call `execvp()` to execute the executable. If the command is not in the background, the shell waits for the last single command to finish using `waitpid()`. Refer to the man pages for function details. The file `ls_grep.c` is an example of process creation and redirection.
- After this, commands like `ls -al` and `ls -al /etc &` should be executable.

##### • File redirection:

- If the command specifies IO redirection files, create them as necessary. Use `dup2()` to change file descriptors (0 for standard input, 1 for output, 2 for error).
- Examples: `ls -al > out`, `cat -q cat 2> dog`, etc.
- Notes:
  - `2>` redirects stderr to the specified file.



- `>&` redirects both stdout and stderr to the specified file.
- `>>` appends stdout to the specified file.
- `>>&` appends both stdout and stderr to the specified file.

- **Pipes:**

- Use `pipe()` to create a pipe for inter-process communication. Redirect the output of one command to the input of the next using `dup2()`. See the example in `ls_grep.c`.
- Commands like `ls -al | grep command` and `ls -al | grep command | grep command.o > out` should work.

- **isatty():**

- When the shell uses a file as standard input, it should not print a prompt. Use `isatty()` to determine if the input is from a file or terminal. This is required for the automated tests to pass.

- **Exit:**

- Implement a special command `exit` that exits the shell when run. It should be parsed by the shell and cause the shell to exit without creating a new process.

#### 4. Submission

- By Monday, September 23 11:59pm:
  - Log in to a CS department machine.
  - Navigate to the `proj3` directory.
  - Run `make clean`.
  - Run `make` to check if the shell builds correctly.
  - Run `make submit_part1`.
- Standard rules apply:
  - Follow the coding standards
  - The shell must compile and
  - Do not look at others' source



## Part 2: Wednesday, October 1, 2014

### Adding Features - Part 2

#### 1. **ctrl-C**

- Make the shell behave like bash when `ctrl-C` is pressed while a command is running, it should interrupt and usually terminate the command. If `ctrl-C` is pressed while the shell is running, it should interrupt and usually terminate the shell. See `crtlc.c` for an example of detecting `ctrl-C`.

#### 2. **Zombie Elimination**

- Set up a signal handler to catch `SIGCHLD` and eliminate zombie child processes. The shell should call `waitpid()` to clean up child processes. The shell should print the message "[PID] exited."

#### 3. **Quotes**

- Add support for quotes in the shell. The shell should support single and double quotes as a single argument. For example, `myshell> ls "command.cc Makefile"` should list `command.cc` and `Makefile` if they exist. The quotes should be removed before using the argument.

#### 4. **Escaping**

- Allow the escape character. Any character after `\` can be part of an argument, including special characters like quotation marks and ampersands. For example, `myshell> echo \"Hello between quotes\"` should print `"Hello between quotes"`, and `myshell> echo this is an ampersand \& this is an ampersand &` should print the correct text with the ampersand.

#### 5. **Builtin Functions**

- Implement the following builtin commands:
  - `printenv` : Prints the environment variables of the shell. The environment variables are stored in `char **environ`, a null-terminated array of strings. Refer to the `environ` man page.
  - `setenv A B` : Sets the environment variable `A` to value `B`.
  - `unsetenv A` : Un-sets environment variable `A`.
  - `source A` : Runs file `A` line-by-line as if it were typed into the shell by a user.
  - `cd A` : Changes the current directory to `A`. If no directory is specified, defaults to the home directory. Refer to the `chdir()` man page.
- These built-ins should be used like other commands, including with redirection and piping.

#### 6. **“.shellrc”**

- When the shell starts, it should attempt to run `“source.shellrc”` equivalent.

#### 7. **Submission**

- By Wednesday, October 2 11:59pm:
  - Log in to a CS department machine.
  - Navigate to the `proj3` directory.
  - Run `make clean`.
  - Run `make` to check if the shell builds correctly.
  - Run `make submit_part2`.
- Standard rules apply as in Part 1.

## Part 3: Friday, October 11 11:59pm

### Moar Features - Part 3

#### 1. Subshells

- Implement subshells. Any argument of the form *(command and args)* will be processed by a child shell process, and the output will be fed back into the original parent shell. For example `(expr 1 + 1)` becomes `echo 2`, and `echo a b > dir; ls $(cat dir)` lists the contents of directories `a` and `b`.
- Steps:
  - Extract the command between `"$("` and `")"` in `shell.l`.
  - Invoke the shell as a child process, passing the command to the child and capturing its output.
  - Place the output obtained from the child into a temporary file for interprocess communication (one to write the command, one to read the output).

#### 2. Environment variable expansion

- Implement environment variable expansion. If `$var` appears in the command, it will be expanded to the value of `var` (or an empty string if `var` is not set). Special variables `$?` (PID of the shell process), `${?}` (run the last command again), `${_}` (last argument in the previous command), and `$_` (last argument in the previous command relative path) should also be expanded. Use `realpath()` to expand the relative path.

#### 3. Tilde expansion

- When `~` appears by itself or before a slash, it should be expanded to the user's home directory. If `~` appears before a word, the expansion should not happen. Tilde expansion should not happen within double quotes.

#### 4. Wildcarding

- Implement wildcarding as follows:
  - First, handle wildcarding on the command line. If a wildcard character (`*` or `?`) is found, insert the file names that match the wildcard into the command line. Use `opendir` and `readdir` to find the entries that match the wildcard (convert from wildcards to regular expressions as needed, see `regular.cc` example).
  - Commands like `echo *`, `echo *.c`, `echo shell.?` should work.
  - Then make it work for any path (examples: `echo */*`, `echo */**/*`). Do not use the `glob()` call.

#### 5. Edit mode

- `tty_raw_mode.c` contains sample code to change the terminal's input from canonical to raw mode. Implement a full line editor in `read_line.c`.
- Add the following code to `shell.l` after the `#include` lines:



```
%{
#include <string.h>
#include "y.tab.h"
////////// Start added code //////////
extern "C" char *read_line();
int mygetc(FILE *f) {
    static char *p;
    char ch;
    if (!isatty(0)) { // stdin is not a tty. Call real getc
        return getc(f);
    }
    // stdin is a tty. Call our read_line.
    if ((p == NULL) || (*p == 0)) {
        p = s; char *s = read_line();
    }
    ch = *p;
    p++;
    return ch;
}
#undef getc
#define getc(f) mygetc(f)
////////// End added code
}%
%%
```

- Modify the Makefile by defining `mygetc` to use `mygetc` in `read_line.c`.
- In `read_line.c`, add the following:
  - Left arrow key: Move the cursor left.
  - Right arrow key: Move the cursor right.
  - Delete key (ctrl-D): Remove the character under the cursor.
  - Backspace key (ctrl-H): Remove the character to the left of the cursor.
  - Home key (ctrl-A): Move the cursor to the beginning of the line.
  - End key (ctrl-E): Move the cursor to the end of the line.
- Do not use the `readline` library.

## 6. History

- Implement a history list. Every time a new line is entered, it should be added to the history list. Implement the following editor commands:
  - Up arrow key: Shows the previous command.
  - Down arrow key: Shows the next command.

## 7. ctrl-R

- Implement the `ctrl-R` feature. It should allow the user to search through the command history by typing part of a command. Pressing "Enter" executes the command.

## 8. Submission

- By Friday, October 11 11:59pm:
  - Add a `README` to the `proj3/` directory with:
    - Features that work as specified in the handout.
    - Features that do not work as specified in the handout.
    - Extra features implemented.
  - Log in to a CS department machine.
  - Navigate to the `proj3` directory.
  - Run `make clean`.
  - Run `make` to check if the shell builds correctly.
  - Run `make submit_part3`.
- Standard rules apply as in Part 1.

# Extra Credit

## 1. Process Substitution (8 points)

- Implement process substitution. Use `mkdtemp()` to generate a temporary directory for a named pipe. Create a named pipe with `mkfifo()` within the temp directory. Connect the substituted command's output to the named pipe and pass the pipe name to the other command. Clean up resources (fifo and directory) after the command is executed using `unlink()` and `rmdir()` functions.

## 2. Path Completion (2 points)

- Implement path completion. When the key is typed, the editor attempts to expand the current word to a matching file similar to Bash.

## 3. Aliases (2 points)

- Implement aliases. Allow users to define commands that encompass other commands (e.g., create shortcuts or add default arguments). For example, an alias for `ls` to always colorize the output, list entries by columns, and append an indicator.

## 4. Jobs (8 points)

- Implement job control features like in Bash:
  - `jobs` : Display the status of jobs in the current session (refer to the `jobs` man page for operands).
  - `fg` : Run jobs in the foreground (refer to the `fg` man page for operands).
  - `bg` : Run jobs in the background (refer to the `bg` man page for operands).
  - `ctrl-Z` : Suspend a foreground process by sending it a `SIGTSTP` signal (implement a similar signal handler as for `ctrl-C`).

# Grading

- Rough breakdown:
  - Part 1: 30 points (tested with `./testall part1`)
  - Part 2: 30 points (tested with `./testall part2`)
  - Part 3: 30 points (tested with `./testall part3`)
  - Grading of readline and `ctrl`

- Penalties:
    - -5 points for memory leaks.
    - -5 points for file descriptor leaks
- The parts build on each other, and

